

Monte Carlo Methods for Estimation: Recap

Mark M. Fredrickson (mfredric@umich.edu)

Module 1: Integrate — Session 3

Computational Methods in Statistics and Data Science (DATASCI 406)

One draw is now a whole dataset

An estimator is a function of the sample, so it is a **random variable**. Its distribution is the **sampling distribution**, and every question we ask about an estimator is an expectation over it.

```
theta_hats <- replicate(k, {  
  x <- rmodel(n, theta)    # simulate a whole dataset of size n  
  estimator(x)             # collapse it to one number  
})
```

- n is the usually given by the a real data set
- k is your our choice: time vs. precision of estimate

Bias, variance, and MSE

$$\text{Bias}(\hat{\theta}) = E(\hat{\theta}) - \theta, \quad \text{MSE}(\hat{\theta}) = E((\hat{\theta} - \theta)^2) = \text{Var}(\hat{\theta}) + \text{Bias}(\hat{\theta})^2$$

- All three are expectation, estimate using mean of $\hat{\theta} - \theta$, or $(\hat{\theta} - \theta)^2$.
- MSE balance accuracy and precision
- The ranking can depend on n , so **try several n**

Coverage and Expected Width

$$\text{coverage} = P(A \leq \theta \leq B) = E(1(A \leq \theta \leq B))$$

- Replicate samples with known θ , create CI for each sample, proportion that include true θ
- Coverage is a proportion out of k : report `binom.test(hits, k)$conf.int` alongside it.
- After establishing coverage, compare different CI methods on **expected width** $E(B - A)$.